

Empirical Noise-Adaptive Circuit-Depth Thresholds on IBM Quantum Heron-Class Devices

George D. Anderson II¹

¹Independent Researcher

January 31, 2026

Abstract

We present an empirical study of circuit-depth thresholds on IBM Quantum Heron-class hardware using 4- and 5-qubit GHZ-state preparation circuits with logical depths from 2 to 20. On `ibm_torino`, we observe qubit-count-dependent depth limits: 4-qubit circuits reach a fidelity plateau near 88% beyond depth 10, while 5-qubit circuits show continued decay from 89.4% to 83.7% over the same range. The observed fidelity at depth 20 substantially exceeds predictions from an independent gate-error model (88% vs. 61%), indicating coherent error effects and partial error cancellation not captured by simple accumulation models. These results provide hardware-specific depth bounds for near-term quantum circuits and support calibration-based depth selection over fixed heuristics. We recommend $\text{max_depth} \leq 10$ for applications requiring $>85\%$ fidelity on this hardware.

Keywords: quantum computing, NISQ devices, circuit depth, noise characterization, quantum hardware, IBM Quantum

1 Introduction

1.1 Motivation

Deeper parameterized quantum circuits (PQCs) express richer functions but accumulate more gate errors—a tradeoff that sits at the heart of near-term quantum algorithm design. At sufficient depth, this tradeoff tips decisively against the practitioner: gradients vanish exponentially in the “barren plateau” regime [1, 6], and accumulated noise degrades measurement outcomes faster than added layers improve expressivity. Knowing exactly *where* that tipping point lies on real hardware is therefore a prerequisite for any principled adaptive circuit strategy.

Despite this importance, most circuit design today relies on one of three imprecise approaches: fixed-depth architectures with arbitrary depth choices, heuristic growth strategies without hardware-validated stopping criteria, or theoretical noise models that do not capture device-specific behavior. The gap between these proxies and measured hardware performance motivates the present work.

1.2 Contributions

We characterize depth-dependent fidelity decay on IBM Quantum’s Heron-class processors across 4- and 5-qubit GHZ circuits at depths 2–20. Concretely, we contribute:

1. **Empirical depth thresholds** on `ibm_torino`: a noise plateau at $d = 10$ for 4-qubit circuits and continuous linear decay for 5-qubit circuits.
2. **Device-specific noise ceiling identification** at 83–88% fidelity, with qubit-count-dependent decay of 3.2% (4-qubit) vs. 5.7% (5-qubit).
3. **Evidence against independent error models**: the observed 88% fidelity at depth 20 substantially exceeds the 61% predicted by independent gate-error accumulation, implicating coherent errors and partial cancellations.
4. **Concrete depth recommendations** for adaptive quantum architectures, including fidelity-stratified thresholds.
5. **An open-source experimental framework** for reproducible hardware characterization.

1.3 Related Work

Barren plateaus. McClean et al. [6] established that gradients of random PQCs vanish exponentially with system size; Cerezo et al. [1] showed this extends to certain cost function choices. Our work provides experimental fidelity measurements that complement these theoretical gradient analyses on real hardware.

Circuit compilation and depth reduction. Work on automated gate synthesis and compilation [2, 7] targets gate count as a proxy for quality. We instead measure the fidelity–depth relationship directly, revealing that gate count is an imperfect predictor of output quality due to coherent error effects.

Hardware noise characterization. Standard benchmarking [3, 5] reports single-gate fidelities and coherence times but does not probe how these compound across entangled multi-qubit circuits at varying depths—the regime most relevant to VQAs.

Adaptive quantum circuits. Proposals for progressive circuit growth [4, 8] assume depth can be increased without bound until a training signal stalls. Our hardware-validated thresholds give these methods a principled stopping criterion they currently lack.

2 Experimental Setup

2.1 Quantum Hardware

- **Platform:** IBM Quantum Platform (cloud access)
- **Processor:** `ibm_torino` (Heron r2 architecture, 133 qubits)
- **Qubit chain:** [55, 65, 66, 67, 68] (selected via error-rate optimization)
- **Native gate set:** CZ (two-qubit), RZ, SX, X (single-qubit)
- **Calibration date:** January 31, 2026

Qubit quality metrics (chain average):

- Single-qubit gate error: $\approx 0.05\%$
- Two-qubit (CZ) gate error: $\approx 0.87\%$
- T_1 (relaxation): $\approx 150 \mu s$
- T_2 (dephasing): $\approx 120 \mu s$
- Readout error: $\approx 1.2\%$

2.2 Circuit Design

We use GHZ state preparation circuits with a depth-parameterized structure chosen to ensure non-trivial entanglement at all depths while preventing the transpiler from cancelling gates across layers.

Base GHZ preparation ($d = 1$):

$$H(q_0) \rightarrow \text{CZ}(q_0, q_1) \rightarrow \cdots \rightarrow \text{CZ}(q_{n-2}, q_{n-1}) \quad (1)$$

Depth extension ($d > 1$): Each additional layer l appends

$$\begin{aligned} &\text{RZ}(\theta_l) \text{ on all qubits} \\ &\text{CZ chain: } (q_i, q_{i+1}) \text{ for all } i \\ &\text{RZ}(-\theta_l) \text{ on all qubits} \end{aligned} \quad (2)$$

where $\theta_l = 0.1 \times l$. Varying rotation angles prevent gate cancellation during transpilation, and CZ count scales predictably as $\approx (n-1) \times 2$ per layer.

2.3 Experimental Protocol

Test configurations:

- System sizes: $n \in \{4, 5\}$ qubits
- Circuit depths: $d \in \{2, 4, 6, 8, 10, 15, 20\}$
- Shots per circuit: 1000
- Total circuits: $2 \times 7 = 14$

Metrics:

1. **GHZ fidelity:** $F = (P(|0 \cdots 0\rangle) + P(|1 \cdots 1\rangle))$
2. **Parity expectation:** $\langle Z^{\otimes n} \rangle$
3. **Heavy output frequency (HOG):** fraction of shots above median probability

Transpilation used Qiskit `optimization_level=1` with a fixed `initial_layout` to ensure consistent physical qubit mapping across all depth configurations.

3 Results

3.1 Fidelity Decay Patterns

Figure 1 shows GHZ fidelity as a function of circuit depth for both system sizes. The two curves tell qualitatively different stories.

The 4-qubit circuit peaks at 92.8% fidelity at $d = 4$, declines to 87.9% by $d = 10$, and then effectively saturates: fidelity at depths 15 and 20 (88.2% and 88.4%, respectively) is indistinguishable within shot noise. Total decay from $d = 2$ to $d = 20$ is only 3.2%. The 5-qubit circuit, by contrast, decays monotonically throughout—from 89.4% at $d = 2$ to 83.7% at $d = 20$, a total loss of 5.7% at a roughly constant rate of 0.3% per depth step. The absence of any plateau in the 5-qubit data suggests that noise accumulates in a qualitatively different regime once a fifth entangled qubit is added.

3.2 Gate Count Scaling

Figure 2 plots fidelity against CZ gate count (left) and gate count against logical depth (right). The right panel confirms the expected linear scaling: 6 CZ gates per depth layer for 4-qubit circuits, 8 per layer for 5-qubit, consistent with the $(n - 1) \times 2$ construction. The left panel is more revealing: the 4-qubit system reaches peak fidelity at just 12 CZ gates ($d = 4$), then degrades before stabilizing. The 5-qubit system never finds a stable operating point, declining steadily across the full gate range. Both systems appear to reach a noise floor beyond 60 gates.

Empirical Depth Thresholds for ADQNN-GAGE ibm_torino Qubit Chain [55, 65, 66, 67, 68]




Figure 1: GHZ fidelity vs. circuit depth for 4-qubit (blue) and 5-qubit (orange) systems on `ibm_torino`. The 4-qubit system saturates near depth 10; the 5-qubit system decays continuously throughout.

3.3 Noise Ceiling and Model Discrepancy

Both system sizes converge to fidelity $> 83\%$ at depth 20, suggesting a device-level noise floor. The likely contributors are cumulative readout error ($\approx 1.2\%$ per qubit, totalling 4.8–6% across the chain), gate errors approaching equilibrium, and thermal relaxation during measurement.

Critically, this floor is far higher than naive error models predict. The independent gate-error model gives:

$$F(d) \approx (1 - \varepsilon)^{N_{\text{gates}}(d)} \quad (3)$$

For $\varepsilon \approx 0.0087$ and $N_{\text{gates}} \approx 6d$ (4-qubit), this predicts $F(20) \approx 61\%$ —27 percentage points below the observed 88%. The gap strongly suggests that coherent error components partially cancel across gates rather than accumulating independently. This is a practically important finding: it means that CZ gate count alone is a poor predictor of circuit output quality, and that the real noise floor on Heron-class hardware is more forgiving than standard error budgeting implies.

Practical implication: For applications requiring $> 85\%$ fidelity, $\text{max_depth} \leq 10$ is recommended on this hardware. At depth 10, the 4-qubit circuit has already saturated near




Figure 2: Left: GHZ fidelity as a function of CZ gate count. Right: CZ gate count scaling with logical circuit depth. Linear fits confirm $6d$ (4-qubit) and $8d$ (5-qubit) scaling.

its noise floor, so further depth adds complexity without improving fidelity.

3.4 Statistical Analysis

Shot noise produces an estimated standard deviation of $\pm 1.2\%$ across repeated measurements. A Student’s t -test confirms $p < 0.05$ for fidelity differences exceeding 2% . Three independent runs at different times of day yield consistent trends, confirming reproducibility within the calibration window.

4 Discussion

4.1 Implications for Quantum Neural Networks

The divergence between 4- and 5-qubit behavior carries a concrete lesson for adaptive quantum circuit design. A fixed-depth strategy set at $d = 15$ would perform reasonably for 4-qubit circuits (fidelity $\approx 88\%$, near the plateau) but wastefully so—those five extra layers beyond depth 10 contribute nothing. For 5-qubit circuits, the same fixed depth pushes the system 0.8% below the 85% threshold, which may not be acceptable for gradient-sensitive training. Adaptive depth that tracks fidelity response is not a convenience; it is the only strategy that reliably lands in the trainable, high-fidelity regime for both system sizes.

A second implication concerns scaling. The 5-qubit circuit degrades faster than simple qubit-count arithmetic would suggest—its decay (5.7%) is nearly twice the 4-qubit decay (3.2%) despite adding only one qubit and two CZ gates per layer. This superlinear sensitivity to system size suggests that entanglement-induced error correlations, not just added gate count, drive the difference. Designers targeting larger systems should anticipate that this trend continues.

4.2 Comparison with Theoretical Predictions

Barren plateau theory predicts exponential gradient vanishing:

$$\text{Var}(\partial_\theta \mathcal{L}) \sim 2^{-n} \quad (4)$$

For $n = 4\text{--}5$ qubits this effect is mild, and our fidelity measurements are consistent with the circuits remaining in the trainable regime throughout—fidelity degrades but does not collapse. Hardware noise, not fundamental gradient vanishing, sets the practical depth limit in this regime.

The more striking comparison is with the independent error accumulation model discussed in Section 3.3. The 27-percentage-point gap between its prediction and observation implies that Heron-class processors benefit from significant coherent error cancellation. Future noise characterization work should investigate whether this cancellation is incidental to GHZ circuits specifically or a broader feature of the CZ gate calibration on this architecture.

4.3 Recommendations for ADQNN-GAGE and Adaptive Architectures

The data support three concrete recommendations for adaptive circuit growth strategies. First, the safe growth zone is $d \in [2, 8]$, where both system sizes maintain $> 87\%$ fidelity and gains from additional depth are still measurable. Second, growth should terminate at $d = 10$: for 4-qubit circuits this is where the plateau begins, and for 5-qubit circuits it marks the point where fidelity crosses the 85% threshold on its way down. Third, any deployment on a new backend should re-run this characterization, since the 88% plateau is specific to `ibm_torino`’s calibration and qubit chain; a different architecture or qubit selection may shift the threshold substantially.

Fidelity-stratified depth limits:

- $> 90\%$ required: `max_depth` = 4–6
- 85–90% tolerated: `max_depth` = 8–10
- $< 85\%$ acceptable: may explore $d > 10$, with the caveat that gains beyond depth 10 are negligible for 4-qubit systems

4.4 Limitations

Several limitations bound the generalizability of these results. The study covers a single device, qubit chain, and entanglement structure; other GHZ topologies or ansatz families may exhibit different decay profiles. With 1000 shots, statistical uncertainty is approximately 3%, which could obscure fine structure near the plateau. Qubit calibrations also drift over days, so these results should be treated as valid within roughly one week of the calibration date. Finally, results on superconducting transmon or trapped-ion architectures would likely differ substantially: trapped-ion gates have different error rates and coherence properties that could shift both the noise floor and the plateau onset.

5 Conclusion

On IBM Quantum’s `ibm_torino`, circuit depth is not a free parameter. Beyond depth 10, a 4-qubit GHZ circuit gains nothing: fidelity saturates at 88% and further layers add only gate overhead. A 5-qubit circuit has no such plateau and crosses the 85% fidelity threshold at depth 10 on its way to 83.7% at depth 20. Optimal depth is therefore hardware- and system-size-specific, not universal.

The gap between the observed 88% noise floor and the 61% predicted by independent error models is the most consequential quantitative result of this study. It suggests that standard error budgeting significantly underestimates the achievable fidelity of entangled circuits on Heron-class hardware, and that coherent error effects merit explicit modeling in future characterization work.

Future experiments should extend these measurements to 10–15 qubit systems (where barren plateaus become the dominant concern), benchmark the same protocol across multiple IBM and IonQ backends, and investigate whether error mitigation techniques can raise the noise floor near the depth threshold.

6 Reproducibility

All code, data, and experimental protocols are available at:

<https://github.com/jeorgexyz/noise-adaptive-depth-probe>

Software versions: Qiskit 1.2.4, qiskit-ibm-runtime 0.15.0, Python 3.11.

Data availability: Raw measurement counts and transpiled circuits are provided in `results/depth_threshold_results.json` in the repository.

References

- [1] Marco Cerezo, Akira Sone, Tyler Volkoff, Lukasz Cincio, and Patrick J Coles. Cost function dependent barren plateaus in shallow parametrized quantum circuits. *Nature Communications*, 12(1):1791, 2021.
- [2] Edward Farhi and Hartmut Neven. Classification with quantum neural networks on near term processors. *arXiv preprint arXiv:1802.06002*, 2018.
- [3] Jay M Gambetta, Jerry M Chow, and Matthias Steffen. Building logical qubits in a superconducting quantum computing system. *npj Quantum Information*, 3(1):2, 2017.
- [4] Edward Grant, Leonard Wossnig, Mateusz Ostaszewski, and Marcello Benedetti. An initialization strategy for addressing barren plateaus in parametrized quantum circuits. *Quantum*, 3:214, 2019.
- [5] Easwar Magesan, Jay M Gambetta, and Joseph Emerson. Scalable and robust randomized benchmarking of quantum processes. *Physical Review Letters*, 106(18):180504, 2011.

- [6] Jarrod R McClean, Sergio Boixo, Vadim N Smelyanskiy, Ryan Babbush, and Hartmut Neven. Barren plateaus in quantum neural network training landscapes. *Nature Communications*, 9(1):4812, 2018.
- [7] Yunseong Nam, Neil J Ross, Yuan Su, Andrew M Childs, and Dmitri Maslov. Automated optimization of large quantum circuits with continuous parameters. *npj Quantum Information*, 4(1):23, 2018.
- [8] Andrea Skolik, Jarrod R McClean, Masoud Mohseni, Patrick van der Smagt, and Martin Leib. Layerwise learning for quantum neural networks. *Quantum Machine Intelligence*, 3(1):5, 2021.

A Detailed Experimental Data

Table 1: 4-Qubit fidelity data (qubits [55, 65, 66, 67]).

Depth	Transpiled Depth	CZ Gates	Fidelity	Std Dev	Job ID (prefix)
2	18	6	0.916	0.009	d5vd9bt...
4	24	12	0.928	0.008	d5vd9g5...
6	30	18	0.908	0.009	d5vd9hi...
8	36	24	0.915	0.009	d5vd9iq...
10	42	30	0.879	0.010	d5vd9pq...
15	57	45	0.882	0.010	d5vd9ra...
20	72	60	0.884	0.010	d5vd9sj...

Table 2: 5-Qubit fidelity data (qubits [55, 65, 66, 67, 68]).

Depth	Transpiled Depth	CZ Gates	Fidelity	Std Dev	Job ID (prefix)
2	22	8	0.894	0.010	d5vdd2i...
4	28	16	0.905	0.009	d5vdd9q...
6	34	24	0.879	0.010	d5vddb3...
8	40	32	0.874	0.011	d5vddcl...
10	46	40	0.863	0.011	d5vddk3...
15	61	60	0.851	0.011	d5vddla...
20	76	80	0.837	0.012	d5vddml...

B Circuit Construction Details

B.1 GHZ Circuit with Depth Parameter

```

def build_ghz_circuit(n_qubits, depth):
    qc = QuantumCircuit(n_qubits, n_qubits)

    # Initial GHZ preparation
    qc.h(0)
    for i in range(n_qubits - 1):
        qc.h(i + 1)
        qc.cz(i, i + 1)
        qc.h(i + 1)

    # Depth extension layers
    for d in range(depth - 1):
        angle = 0.1 * (d + 1)
        for i in range(n_qubits):
            qc.rz(angle, i)
        for i in range(n_qubits - 1):
            qc.cz(i, i + 1)
        for i in range(n_qubits):
            qc.rz(-angle, i)

    qc.measure(range(n_qubits), range(n_qubits))
    return qc

```

Listing 1: GHZ circuit builder with depth extension.

B.2 Qubit Selection Algorithm

```

def find_best_qubit_chain(backend, chain_length=5):
    coupling_map = backend.coupling_map
    G = nx.Graph(coupling_map)
    chains = find_all_paths(G, length=chain_length)
    best_chain = min(chains, key=lambda c: avg_error(c, backend))
    return best_chain
# Result: [55, 65, 66, 67, 68], avg CZ error 0.87%

```

Listing 2: Optimal qubit chain selection.